

ONE-DAY MASTERCLASS

CURTIN EXECUTIVE EDUCATION • DR MICHAEL BORCK

01

THE

Using & delivering AI with judgement.

WU

“If the AI is good at running your work, it's good at running everyone else's. Generic competence is the baseline — not the advantage.

SO WHERE DOES YOUR EDGE LIVE?

TODAY

03

One idea. Two scales.

MORNING

The judgement, at the level of your own work: the trust tool, RTCF, the two-pass demo.

AFTERNOON

The same judgement, lifted to delivery: the five differences, human-in-the-loop, the go/no-go.

The trust tool – taught once, reused at project scale

PART ONE


THE AVERAGE, THE TOOL, YOUR EDGE.

Personal fluency — the judgement that lasts after the next model ships.

WHAT AI ACTUALLY IS

It predicts the next word.


Autocomplete, scaled up. Trained on billions of pages, it learned what tends to follow what. The result looks like understanding. It's pattern-matching.



THE ONE MENTAL MODEL TO KEEP

A convincing average.

Fluent, plausible, usually about-right. But "about right" and "exactly right" are not the same — and knowing which you need is the whole game.



ONCE YOU SEE IT AS AN AVERAGING MACHINE

07

Brilliant here. Quietly wrong there.

BRILLIANT

Drafting, summarising, formatting, brainstorming, explaining, finding patterns — where "the plausible version" is the answer.


UNRELIABLE

Precise facts, figures, citations, multi-step logic, your context, ethics — where there's one right answer it can't verify.

NOT A BUG — A PROPERTY

Confident. Wrong. Same polish.

It never says "I don't know." The fluent-but-wrong answer looks identical to the fluent-and-right one. Nothing in the tone warns you.



The trust tool.

AVERAGE → LEAN IN

Low stakes, "about right" is fine. Let AI run; you glance.

PRECISE → HUMAN IN THE LOOP

High stakes, one right answer. A person owns the decision.

Average or precise? Small or large?

Drafting an internal email. Average is fine, stakes small — lean in.

Board-paper figures. Precision required, stakes large — a human owns every number.

Summarising fifty complaints for themes. Lean in for the themes, then skim for the one complaint the AI smoothed away.

RTCF — two halves.

R · T · F SCAFFOLDING

Weaker & self-hosted models need it spelled out. A frontier model infers it from clear writing.

C ALWAYS MATTERS

Context is the one thing only you supply. No model manufactures your edge.

Unusable. Then sendable.

Without: "Write an email about the project delay."


With RTCF: a senior PM, professional & empathetic, writes a client email explaining a 2-week delay (vendor integration; client values transparency; new date March 15), under 200 words — acknowledge → explain → new date → recommit.

Context is where your edge enters.

THE MONEY MOMENT — LIVE

Add your edge. It gets a soul.

Pass 1: a naive prompt → slick, correct, identical for everyone. Pass 2: layer in what only you hold. Watch the output turn yours.




Redesign one workflow.

Pick a task you own. List its steps. Judge each with the trust tool — where does AI genuinely help, where must your judgement stay in charge? Redesign so AI handles the average and you keep the high-stakes parts. You leave with something to trial next week.

A mental model, a grid, a method.

AI produces a convincing average. The trust tool tells you when to lean in. RTCF directs it well — and Context is the part that survives any model. One question for the afternoon: if this is how I trust AI on a single task, how do I trust it across a whole project?

16



10:30

Morning tea

BACK AT 11:00

PART TWO

WHY AI

DELIVERY IS

how the delivery team has one funded initiative — without the predictable failures.

“I’ve delivered projects before. What’s actually different about an AI project?”

THEY DON'T FAIL MORE OFTEN. THEY FAIL DIFFERENTLY — IN PLACES YOU DIDN'T SEE COMING.

They fail *differently*.

01

"Done" can't be specified.

02

The demo is a trap.

03

The data is the uncertainty.

04

Verification is the product.

05

Generic competence is the baseline.


→

So you design the human in.

DIFFERENCE 01

"Done" won't hold still.


Ask twice, get two answers. There's no single correct output — only a spread. Stop chasing a fixed spec; decide what "good enough" looks like for this use.



DIFFERENCE 02

The demo is a trap.


It shows the happy path. The messy real-world cases are where it falls over. Treat the demo as the start of the hard part — not the end.



DIFFERENCE 03

The data is the uncertainty.


You don't order it; you discover it — its gaps, biases and surprises show up only when you work with it. Run delivery as discovery, and build the learning into the plan.



DIFFERENCE 04

Verification is the product.


Outputs vary and can be confidently wrong. The value lives in how you check them, and where a human stays in the loop. Design the checking deliberately — make it part of what you ship.



DIFFERENCE 05

Generic is the baseline, not the edge.

When every rival runs the same models, the model isn't your advantage — anyone can have it. Compete on the human variation a rival's identical model can't reproduce.



“AI gives you something fluent and roughly-right, fast — then asks you to decide when roughly-right is good enough, and who checks when it isn't.

A LEADERSHIP QUESTION. WHICH IS WHY IT LANDS ON YOUR DESK.

You're the delivery lead.

The board funded one AI initiative. Ship it without the predictable failures. Carry it through three moves — scoping and stress-testing it by interviewing the RetailFlow team: live chatbots with opinions, who disagree, and won't do your job for you. access code: pilot2024

Scope. Verify. Decide.

01

Scope it against reality

Interview Data & CIO. Reconcile "move fast" with "the data isn't ready." → objectives, data needs, a defined "good enough."

02

Stakeholders & human-in-the-loop


Interview MD, frontline, CFO. Decide where a human must stay in charge. → stakeholder plan + checkpoint design.

03

Roadmap & the go/no-go

Build the roadmap with gates and a risk register. Then make the call: Scale, Pivot, or Kill.

28




12:30

Lunch

BACK AT 1:15


THE GO / NO-GO CALL

Scale. Pivot. Kill.



At each gate the question isn't "is it perfect?" — it's whether the evidence justifies continuing. Killing early is success, not failure.

30



2:30
Afternoon tea

BACK AT 3:00

A delivery design, not a theory.

SCOPE

Objectives, data requirements, a defined "good enough."

ROADMAP

Milestones and go/no-go gates.

CHECKPOINTS

A stakeholder plan and human-in-the-loop design.


RISK

A register and a Scale / Pivot / Kill call you can defend.

BACK TO THE PROVOCATION

The edge is the part the tool can't average.

Averages converge. Your taste, your contrarian calls, your specific customers — that's the variation a rival's identical model can't reproduce.



Your action plan.

WEEK

Trial your redesigned workflow. Run one task through the trust tool first.

MONTH

Make RTCF a habit. Start one conversation about where a human stays in the loop.

QUARTER

Scope one real initiative — and defend the Scale / Pivot / Kill call.

Tools that don't expire.

DECISION TOOLS

The trust tool, RTCF, the five differences, Scale / Pivot / Kill — any platform, any model.

A PLAN

Specific commitments for the week, month and quarter.

YOUR WORK

A redesigned workflow and a scoped project design you can act on.

GO DEEPER

michael-borck.github.io/the-human-edge — incl. the two-page voice method.

99

THE AI IS YOUR STARTING POINT, NOT YOUR FINISH LINE

STAY THE ONE

HOLDING THE

michael-borck.github.io/the-human-edge

JUDGEMENT