

The Human Edge

Using and Delivering AI with Judgement

Participant workbook · One-day masterclass
Dr Michael Borck · Curtin Executive Education

The day at a glance

Time	Block
9:00–10:30	The average, the tool, the edge
10:30–11:00	<i>Morning tea</i>
11:00–12:30	Using AI well, yourself
12:30–1:15	<i>Lunch</i>
1:15–2:30	Why AI delivery is different
2:30–3:00	<i>Afternoon tea</i>
3:00–4:00	Designing the human in, then shipping
4:00–4:30	The edge that's left to humans

RetailFlow portal access code: pilot2024 (afternoon sprints only — your facilitator will confirm)

This workbook is also available online at
michael-borck.github.io/the-human-edge

The trust tool

The trust tool: when to trust AI

The most useful thing to understand about today's AI: it produces a convincing **average**. Fluent, well-shaped, usually about right — but “about right” and “exactly right” are not the same, and knowing which you need is the whole game.

The trust tool is two questions.

Q1: average, or precision?

- **Average** = “in the right ballpark” is fine. A first draft, a summary, a starting point you'll shape further.
- **Precise** = a single correct value, where close is the same as wrong. A contract figure, a dosage, a citation, an account balance.

Q2: small, or large?

- **Small** = low stakes. Easy to reverse, cheap to get wrong, no one harmed.
- **Large** = real consequences. Money, safety, reputation, someone's job or wellbeing.

The four corners

	Small (low stakes)	Large (high stakes)
Average is fine	Lean in. Let AI run.	Lean in, with a glance. Sanity-check before it goes out.
Precision required	Use, then verify. Confirm the specifics.	Human in the loop. A person owns the decision.

Two scales

This grid works at two levels — which is why it's the spine of the whole day:

- **At the task** (morning): “*Do I trust this output for my work?*” — average or precise, small or large?
- **At the project** (afternoon): “*Where does a human stay in the loop across this project?*” — the same corners, applied to where verification and oversight are designed in.

The hardest question leaders ask about AI is “where does a human stay in the loop?” Most teams answer it by gut — everywhere (you lose speed) or nowhere (you lose trust). The trust tool makes it *answerable*. You stop arguing about AI in the abstract and start asking, task by task: average or precise, small or large?

Drawn from the companion book, *Conversation, Not Delegation*.

Your prompt library

Your prompt library

RTCF isn't a trick — it's being clear about what you actually want. Use this sheet to build a working set of prompts for tasks you genuinely do. Most prompts need all four only when the work matters; keep it lighter for quick tasks.

The framework

- **R — Role:** who should the AI be? (e.g. “a sceptical senior reviewer”, “a plain-English editor”)
- **T — Task:** what exactly should it do? (one clear verb)
- **C — Context:** what does it need to know? (**this is where your edge enters** — the specifics only you hold)
- **F — Format:** how should the output look? (length, structure, tone)

When does RTCF matter most? The scaffolding (Role, Task, Format) earns its keep on weaker or self-hosted models that need structure spelled out. A frontier model follows clear, natural instructions well — if you brief people well, you're already most of the way there. The part that *always* matters is **Context**: the specifics only you hold. No model supplies that for you.

Your prompts

Write three prompts for real tasks from your role. Start with a naive version, then rebuild it with RTCF.

Prompt 1

Task:

Naive:

RTCF: - R: - **T:**

..... - **C:**

..... - **F:**

.....

Prompt 2

Task:

Naive:

RTCF: - R: - **T:**

..... - **C:**

..... - **F:**

.....

Prompt 3

Task:

Naive:

RTCF: - R: - **T:**

..... - **C:**

..... - **F:**

.....

The two-pass lesson. Run your prompt once with thin context → slick, generic, anyone-could-have-it. Then layer in what only *you* hold (a customer, a contrarian bet, a taste call). Watch the output get a soul. *That's where your edge lives.*

Redesign your workflow

Redesign your workflow

Pick **one task you do repeatedly** — something you own, that takes real time. You'll map it, judge each step against the trust tool, and redesign it so AI handles what it does well while your judgement stays in charge.

1. Name the task

What's the task? (e.g. "weekly stakeholder report", "client onboarding pack", "inquiry triage")

.....

2. List the steps

Break the task into its steps, in order. Aim for 4–8.

Step	What happens	How long
1		
2		
3		
4		
5		

3. Judge each step with the trust tool

For every step: does it need an **average** or **precision**? Are the stakes **small** or **large**?

Step	Average or Precise?	Small or Large?	Lean in / Human in the loop
1			
2			
3			
4			
5			

4. Redesign

Where does AI genuinely help (average, small stakes → lean in)? Where must your judgement stay in charge (precise, large stakes → human in the loop)? Sketch the new version below.

New workflow:

.....

.....

5. The test

Pick **one step** to trial with AI this week. What will you try, and how will you know if it worked?

The point isn't "use AI more." It's *use AI deliberately, exactly where it belongs* — and protect the steps where your judgement is the value.

Customer-service chatbot

Your group is the delivery lead for this funded initiative.

The mandate

The board has approved **\$450K** to deploy an AI customer-service chatbot across RetailFlow's online channels. The case that won funding: customer satisfaction has fallen 78% to 68% and email response times sit at **26 hours**. The chatbot is meant to deflect routine queries (order status, returns, sizing) and cut response times dramatically.

What the board thinks it bought: "a chatbot that answers customers, live in a quarter."

Your job: deliver it without the predictable failures, and define what "good enough to go live" actually means.

Why this one is hard to *deliver* (not just buy)

- It is **customer-facing**, so a confident hallucination isn't an internal embarrassment. It's a wrong returns policy quoted to a real customer, in writing.
- The data looks ready but isn't uniform: returns and sizing rules vary by product line and have exceptions that live in people's heads, not systems.
- "It works in the demo" will be true within week two. That is the trap (Difference 2): the demo handles the easy 80%; the cost is in the angry-customer edge cases.

The data reality (ask Priya)

Conversation/transcript data is **good**, the best of the four initiatives. Priya's honest read: "*good data foundation, ~6-month implementation, proceed with confidence.*" But "good foundation" does not mean "knows your policies." Grounding the bot in current, correct policy documents (RAG) is where the real work is.

Who to interview, and the tension they bring

- **Tom Walsh (Customer Service Manager):** his team gets automated first. The evaluator's-advantage and apprenticeship-pipeline questions land hardest here. Who reviews the bot's answers, and does that role build or consume his juniors' judgement?
- **Priya Sharma (Data):** feasibility and realistic "good enough" thresholds.
- **Emma Rodriguez (Managing Director):** what was promised to the board, and by when.

Your starting question for "good enough"

What deflection rate, and what *maximum acceptable wrong-answer rate*, makes this safe to put in front of customers? Where must a human stay in the loop before a response is sent?

Dynamic pricing

Your group is the delivery lead for this funded initiative.

The mandate

The board has approved **\$850K** to build AI-driven dynamic pricing, automatically adjusting prices across fashion and homewares lines in response to demand, competitor moves, inventory levels and margin targets. The case that won funding: a “significant revenue opportunity” and competitors “already doing it.”

What the board thinks it bought: “smarter prices that lift margin, live this year.”

Your job: deliver it without the predictable failures, and define what “good enough” pricing decisions look like before any price goes live automatically.

Why this one is hard to *deliver*

- Pricing is **precise and high-stakes**, squarely in the danger zone of the trust tool. A wrong automated price is public, instant, and can read as price-gouging or trigger a margin loss across thousands of SKUs.
- It depends on **competitive data RetailFlow doesn’t cleanly have**, plus demand signals split across online and store systems.
- A pricing model that looks great in backtesting (the “demo”) can behave badly in live conditions it never saw. That is exactly the demo-to-production gap, with money attached.

The data reality (ask Priya)

Priya’s honest read: *“high potential but high complexity: needs better competitive data, staff capability gaps. Recommend Phase 2, not Phase 1; 12+ months realistic.”* Expect her to push back on any plan that promises live automated pricing inside a quarter. Your scope has to confront this.

Who to interview, and the tension they bring

- **Priya Sharma (Data):** will tell you the honest timeline; the central tension of your sprint.
- **Marcus Kim (CIO):** wants speed and sees the revenue prize; “competitors aren’t waiting.”
- **David Chen (CFO):** owns the margin numbers and the go/no-go criteria; wants certainty AI can’t give.

Your starting question for “good enough”

Should pricing be **fully automated, or human-approved** above a threshold? What guardrails (price floors/ceilings, max change per cycle) make automated pricing safe, and where must a human approve before a price changes?

Fraud detection

Your group is the delivery lead for this funded initiative.

The mandate

The board has approved **\$650K** to deploy AI fraud detection across RetailFlow's online payments, to cut chargebacks and fraudulent returns.

What the board thinks it bought: "fewer fraud losses, fast."

Your job: deliver it without the predictable failures, and define what "good enough" detection looks like when both false positives and false negatives carry real cost.

Why this one is hard to *deliver*

- The honest delivery question here is **build vs. buy**. Priya's view: a third-party solution is more practical than building in-house. So your "delivery" may be a *vendor integration and oversight* project, not a data-science build. It is a different kind of plan, and a useful contrast for the room.
- The two failure modes pull in opposite directions: too aggressive and you **block legitimate customers** (revenue + reputation hit); too loose and fraud slips through. There is no single "accurate"; it's a tuned trade-off.
- A vendor demo on their data tells you little about performance on *RetailFlow's* fraud patterns (demo-to-production gap, vendor edition).

The data reality (ask Priya)

Priya's honest read: *"third-party AI more practical than building; our data-science team shouldn't own this. Quick win with the right vendor, 3–6 months. Buy, don't build."* Your roadmap should reflect a procurement-and-integration shape, with the threshold tuning as the real risk.

Who to interview, and the tension they bring

- **Priya Sharma (Data)**: argues buy-not-build; what your team should and shouldn't own.
- **David Chen (CFO)**: owns the loss numbers and the ROI case; wants certainty.
- **Tom Walsh (Customer Service)**: his team handles the fallout when a legitimate customer is wrongly blocked. Who reviews flagged transactions, and how fast?

Your starting question for "good enough"

Where do you set the threshold, and **who reviews flagged transactions before a customer is blocked**? How do you measure success when the two errors (false block vs. missed fraud) have very different costs?

Inventory optimisation

Your group is the delivery lead for this funded initiative.

The mandate

The board has approved **\$1.1M** (the largest of the four) to build AI inventory optimisation: forecasting demand and automating replenishment across 50 stores plus the online channel, to cut both stockouts and overstock.

What the board thinks it bought: “major efficiency gains across the whole network.”

Your job: deliver it without the predictable failures, and define a “good enough” forecast that’s safe to act on automatically.

Why this one is hard to *deliver*

- It is the **broadest in scope** (every store, every category), so the temptation to boil the ocean is strong, and the blast radius of a bad forecast is large.
- It runs straight into **legacy systems**: the POS and inventory systems are 10–15 years old, and online/store data is siloed. Integration is where this project lives or dies.
- A forecast that’s “about right” in aggregate (the demo) can be badly wrong for a specific store/category, and that’s where the cost and the staff frustration show up.

The data reality (ask Priya)

Priya’s honest read: “*significant value potential, but data quality issues to resolve first; integration complexity is underestimated. 9–12 months. Start with ONE category pilot, prove value, then scale.*” Her advice basically pre-writes your roadmap, if you ask her.

Who to interview, and the tension they bring

- **Priya Sharma (Data)**: pushes the single-category pilot; data-quality realism.
- **Marcus Kim (CIO)**: knows the legacy-systems pain better than anyone; also impatient for transformation.
- **Sarah Thompson (COO)**: operational reality across stores; the people who’ll act on (or ignore) the forecasts. She slows things down, sometimes rightly.

Your starting question for “good enough”

What forecast accuracy, in which category, justifies **automating replenishment** versus keeping it as a recommendation a store manager approves? How do you scope a pilot narrow enough to prove value but real enough to matter?

Sprint 1: Scope it against reality

Your initiative: _____ Group: _____

You are the delivery lead. The money is committed. Before you can plan anything, you have to scope it against reality, and the only people who know the reality are down the hall. **Interview Priya (Data) and Marcus (CIO)** and use what they tell you.

Not on the table today: whether this initiative is worth the investment. That call is made: the board funded it. This is *delivery* scoping (what ships first, what data it needs, what “good enough” means), not business-case scoping. If your group finds itself arguing about ROI or whether to do it at all, you’ve drifted into a different conversation.

Tip: Marcus will push speed and ambition. Priya will give you the honest, data-grounded timeline. Your scope is where you reconcile the two, and you’ll have to defend it later.

1. Readiness check

Before you scope, a quick reality check across three lenses (see *Delivery frameworks* on the companion site):

- **Data readiness:** do we have it, and is it good enough? Who owns it? What’s missing or messy?

- **Organisational readiness:** who owns this initiative, who’s affected, who’ll resist? Is the sponsor still in? _____
- **Build • buy • host:** how will we deliver it — buy a product, build it, or self-host? _____

Low readiness in any lens doesn’t kill it — it reshapes the scope (start smaller, fix the data first, buy before you build).

2. Scoped objective (one sentence)

What, specifically, will be delivered, and what is deliberately **out** of scope for this first release?

- In scope: _____
- Out of scope (for now): _____

3. Data requirements (from your Priya interview)

What data the project needs	Does RetailFlow have it, and is it good enough?	Work required first

4. Apply the trust tool

Split the work into parts. For each, is it **Average** (AI produces a good-enough answer) or **Precise/high-stakes** (a human must stay in charge)?

Part of the work	Average → automate	Precise/high-stakes → keep a human

5. Define “good enough”

AI output is non-deterministic; there’s no spec to pass. So *you* must define fitness for purpose **before** deployment.

- The single metric that says “safe to go live”: _____
 - The threshold on that metric: _____
 - The **maximum acceptable error rate**, and what kind of error is unacceptable: _____
-

Leave Sprint 1 with: a scoped objective, data requirements, a trust-tool split, and a defined “good enough.” You’ll build the rest of the plan on these.

Sprint 2: Stakeholders & human-in-the-loop

Your initiative: _____ Group: _____

A working AI system isn't a tool. It's a **division of labour between human judgement and machine output**. This sprint you design that division, and the stakeholder plan around it. **Interview Emma (Managing Director), Tom (Customer Service) and David (CFO).**

1. Stakeholder plan

Stakeholder	What they need / fear	How you'll manage them	Power × Interest
Emma (Managing Director / sponsor)			High / High
David (CFO)			
Tom (CS / frontline)			
<i>add others you met</i>			

2. Human-in-the-loop checkpoints

For every point where AI output could be wrong and it matters, design a checkpoint.

Checkpoint (where a human checks AI output)	Who has the expertise to actually evaluate it?	Does this checkpoint build that expertise, or just consume it?

The evaluator's advantage: experienced people spot the confident-but-wrong answer because they know what right looks like. Judgement is unevenly distributed, so staff your checkpoints with the *right* humans, not just any human.

The apprenticeship trap: if you automate all the grunt work, you stop producing the next generation of evaluators. Mark at least one checkpoint that you keep human *deliberately, to build capability*, not only for quality.

3. The one you'll defend

Which single human-in-the-loop decision are you most confident is non-negotiable, and why?

Leave Sprint 2 with: a stakeholder plan and a human-in-the-loop checkpoint design where each checkpoint names its evaluator and is marked builds-vs-consumes.

Sprint 3: Roadmap, risk & the go/no-go

Your initiative: _____ Group: _____

Now turn your scope, stakeholders and checkpoints into a delivery plan, then defend it when the staff push back and a crisis hits. Remember Difference 2: **progress is illusory**, so your gates exist to catch the demo-to-production gap before it reaches a customer or the board.

1. Delivery roadmap with go/no-go gates

Phase	What happens	Go/no-go gate: what must be true to proceed?
Pilot (narrow)		
Evaluate		
Scale / Pivot / Kill		
Production		

2. Risk register

Risk	AI-specific? (non-determinism / data / demo-gap / verification)	Likelihood	Impact	Mitigation

3. Crisis response (DDCD)

Your facilitator will hand your group a crisis. Respond:

- **Diagnose:** is this technical, people, leadership, or ethical? _____
- **Decide:** your call under uncertainty: _____
- **Communicate:** who, and what do you say? _____
- **Document:** the decision and the reason: _____

4. Assumptions register

Every plan rests on assumptions. Name the ones your plan depends on, and how sensitive it is to each (see *Delivery frameworks* on the companion site). This is what makes the gate-1 call defensible.

Assumption (what must be true)	Sensitivity (if it's wrong, the plan...)	Status at the gate
		holding / shaky / broken
		holding / shaky / broken
		holding / shaky / broken

Scale = load-bearing assumptions holding. **Pivot** = one shaky but fixable. **Kill** = a core assumption has broken.

5. The gate-1 decision

At your first evaluation gate, the data is ambiguous. Make the call and defend it with evidence:

SCALE / PIVOT / KILL (circle one), because: _____

Killing a project that isn't working, early, is success, not failure. The cost of wrongly scaling is far higher than the cost of wrongly killing.

Leave Sprint 3 with: a roadmap with gates, a risk register, a documented crisis response, and a defended go/no-go decision. Together with Sprints 1–2, that's your delivery design.